

Bio Resources at CARC

Center for Advanced Research Computing
University of Southern California
Tomasz Osinski, PhD
Research Facilitator in Life Science



USC

Advanced Research Computing
Enabling scientific breakthroughs at scale

Bio Resources (updated twice a year)

<https://carc.usc.edu/user-information/bio-resources>

The screenshot shows the USC Advanced Research Computing (CARC) website. At the top, a yellow navigation bar contains the text "USC | USC ITS | Office of the CIO | Office of Research" and a search input field labeled "ENTER:". Below this is the USC logo and the text "Advanced Research Computing | Enabling scientific breakthroughs at scale". The main navigation menu includes "About", "Services", "User Information" (circled in red), "Education & Outreach", "News & Events", and "User Support" (highlighted in yellow). The "User Information" dropdown menu is open, showing links: "Getting Started", "Accounts And Allocations", "System Information", "User Guides", "CARC User Portal", "CARC User Forum", "CARC OnDemand", "Condo Cluster Program", "Bio Resources" (circled in red), and "Frequently Asked Questions". The "Bio Resources" link is highlighted in blue. Below the navigation menu is a large banner with a background image of a field. The banner text reads "CARC OnDemand Service Now Available" and "New online access point provides users with web access to CARC file systems and clusters". A "Read More" button is located at the bottom of the banner.

Bio Resources (updated twice a year)

<https://carc.usc.edu/user-information/bio-resources>

- **Genomes** - reference sequences and annotations for commonly analyzed organisms
- **Genbank** - collection of all public nucleotide sequences and their protein translations
- **Genome Taxonomy Database (GTDB)** - an initiative to establish a standardized microbial taxonomy based on genome phylogeny
- **Pfam Database** - large collection of protein families, each represented by multiple sequence alignments and hidden Markov models (HMMs)
- **TIGRFAMs** - curated multiple sequence alignments, Hidden Markov Models (HMMs) for protein sequence classification
- **UniProt** – UniProtKB (curated protein information), UniRef (closely related sequences), UniParc (all protein sequences, consisting only of unique identifiers and sequences)

Bio Resources (through command line)

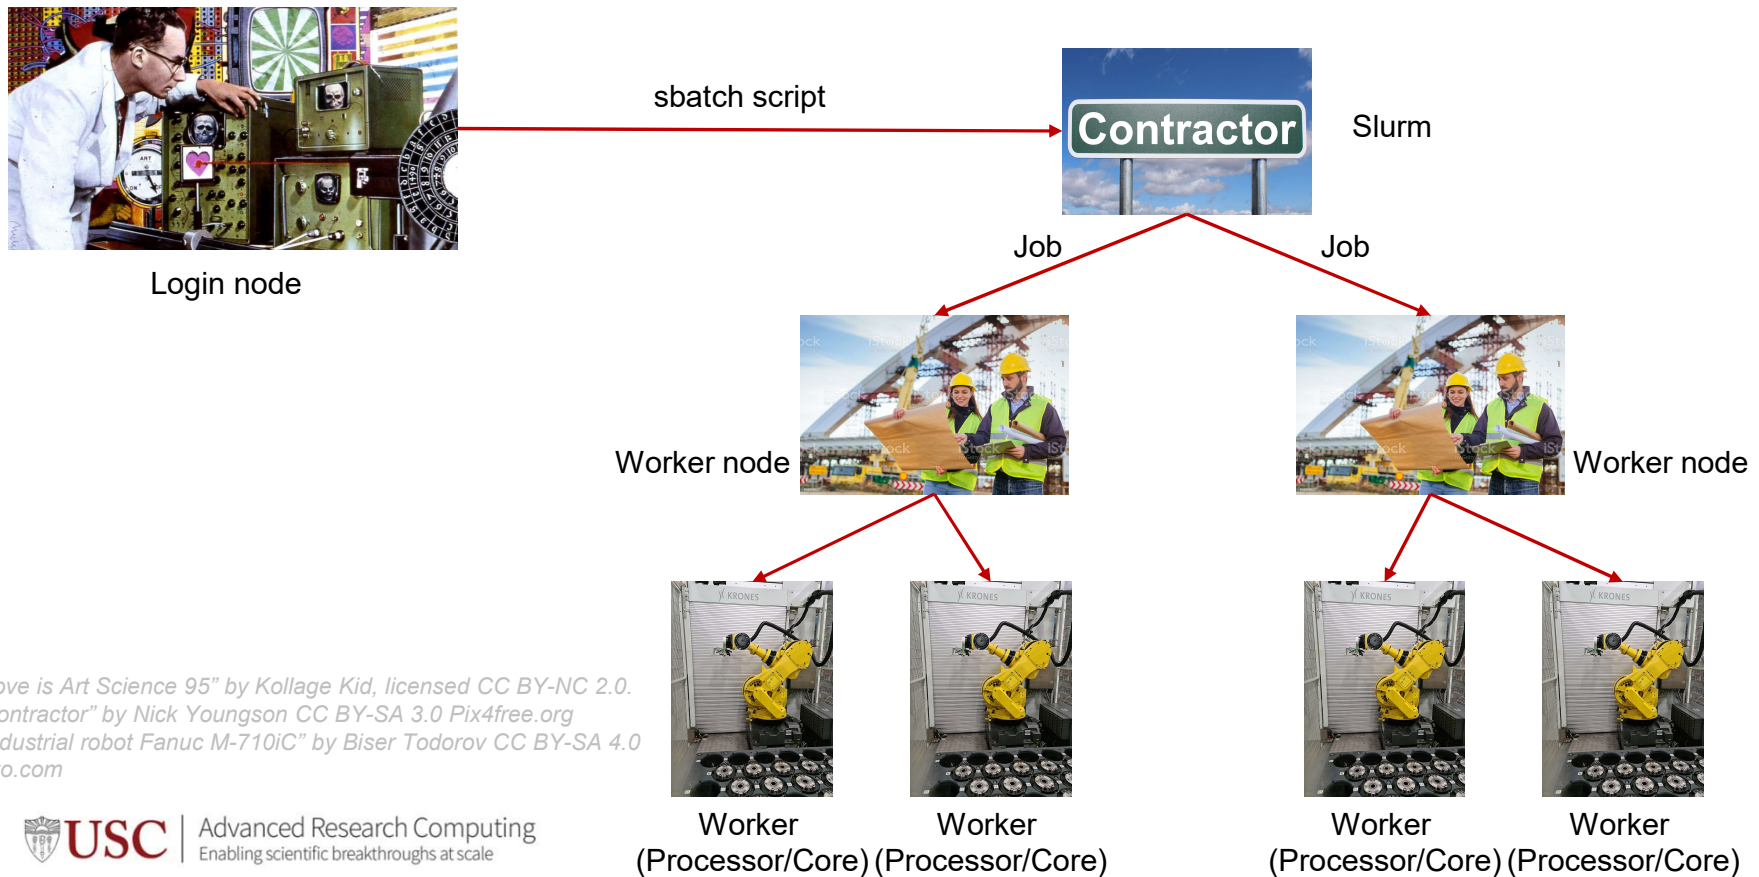
<https://carc.usc.edu/user-information/bio-resources>

- **Biogeotrases** - set of metagenomes, collected under the auspices of the bioGEOTRACES component of the international GEOTRACES program
- **TaraOceans** - marine microbial metagenomes sampled across space and time
- **Variant Effect Predictor cache** - VEP can use a variety of annotation sources to retrieve the transcript models used to predict consequence types. Cache contains all transcript models, regulatory features and variant data for a species and allows for an offline use of VEP

Some terms

- **Head Node** – The system that controls the cluster
- **Worker (Compute) Node** – Systems that perform the computations in a cluster
- **Login Node** – System that users log into to use a cluster
- **Scheduler** – Software that controls when jobs are run and the node they are run on
- **Shell** – A program that users employ to type commands
- **Script** – A file that contains a series of commands that are executed
- **Job** – A chunk of work that has been submitted to the cluster

How does it work?



What partition should I use?

- **debug** - small, short or test jobs that take less than 30m; short queue
- **main** (default) - most jobs; up to 48h runtime (2 days)
- **epyc-64** - larger multithreaded jobs; up to 48h (2 days)
- **largemem** - jobs that require huge amount of memory (up to 1TB); up to 168h (7 days)
- **oneweek** - long running jobs; up to 168h (7 days)

Log into CARC

- Open the terminal:
 - Mac: Applications>Utilities>Terminal or open Spotlight and start typing “terminal”
 - Windows: Start menu>cmd (or use PuTTY or Cygwin)
 - Linux: System tools>Terminal or Accessories>Terminal or search for Terminal
- Type `ssh ttrojan@discovery1.usc.edu`
- Enter your password
- Choose an option in Duo-2FA, and confirm your access
- Answer “No” when asked to save your password
- If successful, your prompt should look something like:
[[ttrojan@discovery1](https://ttrojan@discovery1.usc.edu) ~]

Review of Important Commands: squeue

- Shows the status of jobs running in the queues

```
[osinski@discovery1 ~]$ squeue | head
4679566      main discover sunwool  R    2:19:33      8 d23-[13,15-16],e21-14,e22-[08-09,12],e23-01
4680126      main discover sunwool  R    39:11       8 d23-[13-16],e22-[05-06,08-09]
4678655      main job.slur liukuang  R   11:09:20     1 d14-08
4679445      main 1086-7B asareh   R    4:18:00     1 d11-46
4679444      main 1086-7B asareh   R    4:19:31     1 d05-40
```

- You can also show the status of just your jobs

```
[osinski@discovery1 ~]$ squeue -u ttrojan
      JOBID PARTITION   NAME     USER ST       TIME  NODES NODELIST(REASON)
3678639    epyc-64   test_1  ttrojan PD        0:00      4 (Resources)
3678721    epyc-64   test_2  ttrojan PD        0:00      4 (Priority)
3675759    epyc-64   test_3  ttrojan R   1-01:48:12     2 b22-[29-30]
```

- CD is completed, R is running, PD is waiting to run (pending)

Review of Important Commands: sinfo

- Show the properties of queues and nodes

```
[osinski@discovery1 ~]$ sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up       30:00      6   idle a02-26,e05-[42,76,78,80],e22-13
epyc-64    up    2-00:00:00     32  alloc b22-[01-32]
main*      up    2-00:00:00     16 alloc$ d11-[02-04],e16-[01-13]
main*      up    2-00:00:00     11  maint e16-[14-24]
main*      up    2-00:00:00      4 drain* d17-[03,30],d23-[11-12]
main*      up    2-00:00:00      6  down* d17-[39-44]
main*      up    2-00:00:00      1  drain d06-23
oneweek    up    7-00:00:00      4   mix e01-[52,76],e02-[70-71]
oneweek    up    7-00:00:00      9  alloc e01-[46,60],e02-[40-46]
oneweek    up    7-00:00:00     34  idle e01-[48,62,64],e02-[48-69,72-80]
largemem   up    7-00:00:00      2   mix a16-[02,04]
largemem   up    7-00:00:00      1  alloc a16-03
```

Review of Important Commands: sbatch

- Submit a job to the cluster
- `--partition` partition you want to submit to
 - Default is “main”
- `--nodes` Number of nodes
 - Default is 1
- `--ntasks` Number of CPUs per node
 - Default is 1
- Many more options
 - <https://slurm.schedmd.com/sbatch>

Review of Important Commands: module

- Loads the necessary environment for a program
- `module avail`
 - Shows all modules available, or all the software installed
- `module load`
 - Load the environment for a program
- `module list`
 - Shows modules loaded
- `module unload`
 - Removes a loaded module
- `module purge`
 - Removes all loaded modules

Pay attention to module messages

```
[ttrojan@discovery1 ~]$  
[ttrojan@discovery1 ~]$ module purge  
[ttrojan@discovery1 ~]$ module load gcc/9.2.0  
[ttrojan@discovery1 ~]$ module load bowtie2  
[ttrojan@discovery1 ~]$ module load bedtools2  
[ttrojan@discovery1 ~]$ module load fastqc  
[ttrojan@discovery1 ~]$ module load blast-plus
```

Review: Transferring Files

- SFTP
 - Cyberduck (OSX, Windows) <https://cyberduck.io>
 - WinSCP (Windows) <https://winscp.net/eng/index.php>
 - FileZilla (OSX, Windows, Linux) <https://filezilla-project.org>
- Globus Online (Best way to get data from CARC)
 - go to <https://www.globus.org> in your browser and click **Log In**
 - Search for **University of Southern California** in the box that says "Use your existing organizational login"
- Command line tools - `rsync` or `scp`

More info:

<https://carc.usc.edu/user-information/user-guides/data-management/transfer-overview>

Resources

- CARC home page
 - <https://carc.usc.edu>
- Bio Resources at CARC
 - <https://carc.usc.edu/user-information/bio-resources>
- CARC User Forum
 - <https://hpc-discourse.usc.edu/categories>
- SLURM tutorials
 - <https://slurm.schedmd.com/tutorials.html>
- SLURM quick reference
 - <https://slurm.schedmd.com/pdfs/summary.pdf>

Resources

- CARC home page
 - <https://carc.usc.edu>
- Bio Resources at CARC
 - <https://carc.usc.edu/user-information/bio-resources>
- CARC User Forum ← the most value for the community!
 - <https://hpc-discourse.usc.edu/categories>
- SLURM tutorials
 - <https://slurm.schedmd.com/tutorials.html>
- SLURM quick reference
 - <https://slurm.schedmd.com/pdfs/summary.pdf>

Review: Interactive Jobs

- When you need to provide unpredictable input

```
[ttrojan@discovery1 ~]$ hostname
discovery1.usc.edu
[ttrojan@discovery1 ~]$ salloc -p debug --ntasks=1 -mem=4g
[ttrojan@a02-26 ~]$ hostname
a02-26.hpc.usc.edu
[ttrojan@a02-26 ~]$ exit
exit
[ttrojan@discovery1 ~]$ hostname
discovery1.usc.edu
[ttrojan@discovery1 ~]$
```

Review: Bash Scripts

- Bash scripts are a series of commands that can be grouped together within files to accomplish a series of tasks
- This allows you to run one command instead of several successive commands

Exercise:

- Start an interactive job to the debug queue
- This program sleeps for 10 seconds and then prints out “Hello World”
- Make this file, give it execute permissions, and run

```
#!/bin/bash
# This program: sleeps for 10 seconds, then prints "Hello World"
sleep 10
echo "Hello World"
```

Bash Variables

```
cd raw-seq  
i=1  
ls -l yeast_${i}_50K.fastq  
i=2  
ls -l yeast_${i}_50K.fastq
```

Lets get going

- Detailed policies and directions
 - <https://carc.usc.edu/user-information/getting-started>
- Do not install software yourself, contact us
 - <https://carc.usc.edu/education-and-outreach/office-hours> (Tue, 2:30-5:00)
 - Submit a ticket! (<https://carc.usc.edu/user-support/>)
 - When we install software, it is available to everyone
- Program running slow? *Submit a ticket!*
- Don't know what resources to use? *Submit a ticket!*
- Any other questions? *Submit a ticket or visit our forum*

Prepare to Run Jobs

- Copy example data to your home directory

```
[ttrojan@discovery1 ~]$  
[ttrojan@discovery1 ~]$ git clone https://github.com/uschpc/workshop-bioresources.git  
[ttrojan@discovery1 ~]$ cd workshop-bioresources  
[ttrojan@discovery1 ~]$ ls
```

Create the FastQC Job Script

- Use a text editor to create a file name samplefastqc.sh that contains what follows:

```
#!/bin/bash
#SBATCH --nodes 1
#SBATCH --ntasks 1
#SBATCH --partition debug
#SBATCH --account=ttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
module purge
module load gcc/9.2.0
module load fastqc
echo "Example FastQC start"
sleep 20
fastqc -o results/fastqc-rawseq raw-seq/yeast_1_50K.fastq
echo "Example FastQC end"
```

Run the FastQC Job Script

- Submit the job

```
[ttrojan@discovery1 ~]$ sbatch fastqc1.job  
Submitted batch job 33723
```

- Check the status of the job

```
[osinski@discovery1 ~]$ squeue -u osinski  
JOBID PARTITION      NAME      USER ST      TIME  NODES NODELIST(REASON)  
33723 debug fastqc.s osinski  R       0:02      1 a02-26
```

Check Output File for Errors

- Check Output File for Errors

```
[ttrojan@discovery1 ~]$ cat slurm-33723.out
Started analysis of yeast_1_50K.fastq
Approx 5% complete for yeast_1_50K.fastq
Approx 10% complete for yeast_1_50K.fastq
Approx 15% complete for yeast_1_50K.fastq
Approx 20% complete for yeast_1_50K.fastq
Approx 25% complete for yeast_1_50K.fastq
Approx 30% complete for yeast_1_50K.fastq
Approx 35% complete for yeast_1_50K.fastq
Approx 40% complete for yeast_1_50K.fastq
Approx 45% complete for yeast_1_50K.fastq
Approx 50% complete for yeast_1_50K.fastq
Approx 55% complete for yeast_1_50K.fastq
Approx 60% complete for yeast_1_50K.fastq
Approx 65% complete for yeast_1_50K.fastq
Approx 70% complete for yeast_1_50K.fastq
Approx 75% complete for yeast_1_50K.fastq
Approx 80% complete for yeast_1_50K.fastq
Approx 85% complete for yeast_1_50K.fastq
Approx 90% complete for yeast_1_50K.fastq
Approx 95% complete for yeast_1_50K.fastq
Approx 100% complete for yeast_1_50K.fastq
Analysis complete for yeast_1_50K.fastq
```


Important Things to Note

- Job length
 - If over 24 hours, can this be split up, can threads be increased?
- Many small files
 - To be avoided!
 - Group into larger files
- Data
 - Save space by removing temp files
 - Archive data as soon as reasonable
 - Let us know if you are adding several TB of data
 - Use /scratch or /scratch2 whenever possible for temporary files

Important Things to Note

- Make sure you are not on the login node when you launch an application
 - You can check the system you are on by typing `hostname`
- Make sure you reserve as many processors as you need
 - A mismatch here can increase your runtime or wait time
- Make sure you reserve as much RAM as needed
 - Overestimating increases wait time, underestimating crashes
- Know which resources work the best
 - Sometimes using a debug or epyc-64 is better

Important Things to Note

**No HIPAA Data is allowed on
the cluster!**

(we are working on that part)

SLURM Environment Variables

- `$SLURM_JOB_ID` – The job number
 - Assigned automatically by SLURM
- `$SLURM_JOB_NAME` – The name of the job
 - Similar to the script name
 - Or you can specify one with `-J`
- `$SLURM_NTASKS` – Number of reserved Processors
 - Assigned automatically by SLURM
 - Value is the multiple of the `--ntasks` and `--nodes` values

Multi-Processor Jobs

- The program must support it!
- Our default nodes have mostly 20 cores. Some programs loose efficiency after 8 or 16 processors.
- Wait time and run time adds up if not properly submitted
- Try “program --help” or “man program”
- Use `$SLURM_NTASKS`

Create the BLAST Job Script

- Replace **swissprot** with the path to the v5 of swissprot db obtained from <https://carc.usc.edu/user-information/bio-resources/genbank>

```
#!/bin/bash
#SBATCH --nodes 1
#SBATCH --ntasks 10
#SBATCH --partition debug
#SBATCH --time 00:05:00
#SBATCH --account=tttrojan_001
#SBATCH --chdir /home1/tttrojan/workshop-bioresources
module purge
module load gcc/9.2.0
module load blast-plus
echo "Start BLAST Job"
blastp -db swissprot -query blast/query.txt -out results/blast/results.txt -num_threads
$SLURM_NTASKS
echo "Finish BLAST Job"
```

Run the BLAST Job Script

- Submit the job

```
[ttrojan@discovery1 ~]$ sbatch blast1.job  
Submitted batch job 4773117
```

- Check the status of the job

```
[ttrojan@discovery1 ~]$ squeue -u ttrojan
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
4773117	Main	blast1.j	ttrojan	R	0:02	1	a02-d11

Check BLAST Job Stats with sacct

- sacct can get stats for a job after its completed

<https://slurm.schedmd.com/sacct.html>

```
[ttrojan@discovery1 ~]$ sacct -j 4773117 --format=JobID,State,Elapsed,NCPUS,MaxRSS
```

```
[ttrojan@discovery1 ~]$ sacct -j 4773117 --format=JobID,State,Elapsed,NCPUS,MaxRSS
```

JobID	State	Elapsed	NCPUS	MaxRSS
4773117	COMPLETED	00:00:09	10	
4773117.bat+	COMPLETED	00:00:09	10	1228K
4773117.ext+	COMPLETED	00:00:09	10	832K

GPU Jobs – Example

Use gpu partition

Reserve gpus with --gres parameter

```
#!/bin/bash
#SBATCH --nodes 1
#SBATCH --ntasks 1
#SBATCH --gres=gpu:p100:1
#SBATCH --mail-user=ttrojan@usc.edu
#SBATCH --mail-type=ALL
#SBATCH --chdir /home1/ttrojan
#SBATCH --account=ttrojan_001
module load gcc/8.3.0
module load cuda/10.0.130
module load motioncor2
```

Job Arrays

- A way to run the same commands on many (hundreds, thousands) of datasets/samples.
- A variable called `$SLURM_ARRAY_TASK_ID` is used to determine the element of the array being run.
- `#SBATCH --array=1-1000`
- `$SLURM_ARRAY_TASK_ID` becomes 1 in first job, 2 in second job, etc...

Without Job Arrays – Numbered Files

```
#!/bin/bash
#SBATCH --nodes 1
#SBATCH --ntasks 1
#SBATCH --partition main
#SBATCH --time 00:05:00
#SBATCH --account=ttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
module purge
module load gcc/9.2.0
module load fastqc
echo "Starting FastQC job"
fastqc -o results/fastqc-rawseq-ordered raw-seq-ordered/yeast_1_50K.fastq
fastqc -o results/fastqc-rawseq-ordered raw-seq-ordered/yeast_2_50K.fastq
fastqc -o results/fastqc-rawseq-ordered raw-seq-ordered/yeast_3_50K.fastq
fastqc -o results/fastqc-rawseq-ordered raw-seq-ordered/yeast_4_50K.fastq
fastqc -o results/fastqc-rawseq-ordered raw-seq-ordered/yeast_5_50K.fastq
fastqc -o results/fastqc-rawseq-ordered raw-seq-ordered/yeast_6_50K.fastq
echo "Finish FastQC job"
```

Job Arrays – Numbered Files

- Here is an example SLURM script for a job array. Save as `fastqc_numbered_array.job`

```
#!/bin/bash
#SBATCH --nodes 1
#SBATCH --ntasks 1
#SBATCH --partition main
#SBATCH --time 00:05:00
#SBATCH --array=1-6
#SBATCH --account=ttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
module purge
module load gcc/9.2.0
module load fastqc
echo "Starting FastQC job"
sleep 20
fastqc -o results/fastqc-rawseq-ordered-arr raw-seq-
ordered/yeast_${SLURM_ARRAY_TASK_ID}_50K.fastq
echo "Finish FastQC job"
```

View Job Array

squeue -u uscnetid

```
[ttrojan@discovery1 bio-bootcamp]$ squeue -u ttrojan
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
1152	main	bash ttrojan	R	2:17:32	1	d05-40	
1153	main	bash ttrojan	R	2:17:12	1	d05-40	
1207_1	main	numbered ttrojan	R	0:02	1	d05-41	
1207_2	main	numbered ttrojan	R	0:02	1	d05-40	
1207_3	main	numbered ttrojan	R	0:02	1	d05-42	
1207_4	main	numbered ttrojan	R	0:02	1	d05-45	
1207_5	main	numbered ttrojan	R	0:02	1	d05-44	
1207_6	main	numbered ttrojan	R	0:02	1	d05-44	

Job Arrays – Unnumbered Files

- Start by creating a list of all of the unnumbered filenames
- Then create slurm array script for fastqc jobs that have unnumbered filenames

```
#!/bin/bash
#SBATCH --nodes 1
#SBATCH --ntasks 1
#SBATCH --partition main
#SBATCH --time 00:05:00
#SBATCH --array=1-6
#SBATCH --account=ttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
module purge
module load gcc/9.2.0
module load fastqc echo "Starting FastQC job"
sleep 20
ls raw-seq/ > unnumbered-filenames.txt
line=$(sed -n -e "$SLURM_ARRAY_TASK_ID p" unnumbered-filenames.txt)
fastqc -o results/fastqc-rawseq-unordered raw-seq/${line}
echo "Finish FastQC job"
```

Job Dependencies

- Instructions on how jobs relate to other jobs
- Useful for if you want to run a series of jobs that depend on the output from other jobs
- Examples:

`-d depend=afterok:jobid`

Starts after jobid has finished without errors.

`-d depend=afterok:jobid,before:jobid2`

Starts after jobid is finished, but not until jobid2 has started.

`-d depend:afterok:jobid`

`-d depend:afterok:jobid2`

Starts after both jobid and jobid2 have finished.

`-d depend=afterokarray:jobid`

Starts after the job array jobid has finished without errors.

Job Dependencies

- Why would you do this:
 - Mostly for job pipelines, a series of programs that depend on each other's output that are all submitted at once.

Example:

Step 1:

```
[ttrojan@discovery1 ~]$ sbatch preprocessing-step.sh  
Submitted batch job 18866
```

Step 2:

```
[ttrojan@discovery1 ~]$ sbatch -d after:18866 job-array-step.sh  
Submitted batch job 18870
```

Step 3:

```
[ttrojan@discovery1 ~]$ sbatch -d afterok:18870 postprocessing-step.sh  
Submitted batch job 18867
```


What is Wrong

```
#!/bin/bash
# -----SLURM Parameters-----
#SBATCH --partition main
#SBATCH --ntasks 1
#SBATCH --mem=1g
#SBATCH --nodes 1
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
# -----Load Modules-----
# -----Commands-----
python3 /home1/ttrojan/script.py
```

What is Wrong

- The module is not loaded

```
#!/bin/bash
# -----SLURM Parameters-----
#SBATCH --partition main
#SBATCH --ntasks 1
#SBATCH --mem=1g
#SBATCH --nodes 1
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
# -----Load Modules-----
module purge
module load gcc/9.2.0
module load python/3.7.6
# -----Commands-----
python3 /home1/ttrojan/script.py
```

What is Wrong II

```
#!/bin/bash
# -----SLURM Parameters-----
#SBATCH --partition main
#SBATCH --ntasks 20
#SBATCH --mem=10g
#SBATCH --nodes 1
# -----Load Modules-----
module purge
module load gcc/9.2.0
module load blast-plus
# -----Commands-----
blastn -query fasta.file -db database_name -outfmt 6 \
-num_alignments 1 -num_descriptions 1 -out output_file
```

What is Wrong II

- Number of processors and no working directory

```
#!/bin/bash
# -----SLURM Parameters-----
#SBATCH --partition main
#SBATCH --ntasks 20
#SBATCH --mem=10g
#SBATCH --nodes 1
#SBATCH --account=ttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
# -----Load Modules-----
module purge
module load gcc/9.2.0
module load blast-plus
# -----Commands-----
blastn -query fasta.file -db database_name -outfmt 6 num_alignments 1 \ -num_descriptions 1 -out
output_file -num_threads 20
```

What is Wrong II

- Number of processors and no working directory
- Better to use `$SLURM_NTASKS`

```
#!/bin/bash
# -----SLURM Parameters-----
#SBATCH --partition main
#SBATCH --ntasks 20
#SBATCH --mem=10g
#SBATCH --nodes 1
#SBATCH --account=ttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
# -----Load Modules-----
module purge
module load gcc/9.2.0
module load blast-plus
# -----Commands-----
blastn -query fasta.file -db database_name -outfmt 6 num_alignments 1 \ -num_descriptions 1 -out
output_file -num_threads $SLURM_NTASKS
```

What is Wrong III

```
#!/bin/bash
# -----SLURM Parameters-----
#SBATCH --partition main
#SBATCH --ntasks 1
#SBATCH --mem=200g
#SBATCH --nodes 1
#SBATCH --account=ttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
# -----Load Modules-----
module purge
module load gcc/8.3.0
module load R
# -----Commands-----
Rscript /home1/ttrojan/R_example.R
```

What is Wrong III

- Wrong partition/mem requirements too high

```
#!/bin/bash
# -----SLURM Parameters-----
#SBATCH --partition epyc-64
#SBATCH --ntasks 1
#SBATCH --mem=200g
#SBATCH --nodes 1
#SBATCH --account=ttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
# -----Load Modules-----
module purge
module load gcc/8.3.0
module load R
# -----Commands-----
Rscript /home1/ttrojan/R_example.R
```

What is Wrong IV

```
#!/bin/bash
# -----SLURM Parameters-----
#SBATCH --partition main
#SBATCH --nodes 1
#SBATCH --mem=4g
#SBATCH --ntasks 1
#SBATCH --account=ttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
# -----Load Modules-----
module purge
module load gcc/8.3.0
Module load cuda/10.0.130
module load motioncor2
# -----Commands-----
python /home1/ttrojan/motioncor2.job
```


What is Wrong IV

- GPU resources not specified

```
#!/bin/bash
# -----SLURM Parameters-----
#SBATCH --partition main
#SBATCH --nodes 1
#SBATCH --mem=4g
#SBATCH --ntasks 1
#SBATCH --gres=gpu:p100:1
#SBATCH --account=tttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
# -----Load Modules-----
module purge
module load gcc/8.3.0
Module load cuda/10.0.130
module load motioncor2
# -----Commands-----
python /home1/ttrojan/motioncor2.job
```

What is Wrong V

```
# -----SLURM Parameters-----  
#SBATCH --partition main  
#SBATCH --ntasks 1  
#SBATCH --mem=15g  
#SBATCH --nodes 1  
#SBATCH --account=ttrojan_001  
#SBATCH --chdir /home1/ttrojan/workshop-bioresources  
# -----Load Modules-----  
module purge  
module load gcc/9.2.0  
module load samtools  
# -----Commands-----  
samtools stats example.bam
```

What is Wrong V

- No bash shebang line, `#!/bin/bash`
- Can use long names for SBATCH parameters

```
#!/bin/bash
# -----SLURM Parameters-----
#SBATCH --partition main
#SBATCH --ntasks 1
#SBATCH --mem=15g
#SBATCH --nodes 1
#SBATCH --account=ttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
# -----Load Modules-----
module purge
module load gcc/9.2.0
module load samtools
# -----Commands-----
samtools stats example.bam
```

Genome mapping and tools: Read mapping

- Aim: to find coordinates of reads in the reference genome.
- Challenges:
 - Millions of short sequences
 - Sequences are often paired
 - Errors are not randomly distributed
- Most popular programs are bowtie and bwa (both use Burrows-Wheeler Transform algorithm). Two-step approach:
 - Create an index for the reference genome (one time for one genome).
 - Map reads to the reference genome using this index

Genome mapping and tools – overview I

- FastQC
 - FastQC is a quality control application for high throughput sequence data
 - Checks the quality of their sequence data
 - Generates an HTML report

<http://www.bioinformatics.babraham.ac.uk/projects/fastqc/>

Genome mapping and tools – overview II

- bowtie
 - The first version of bowtie [Langmead et al. 2009] is optimal for:
 - short reads (under 50 bp)
 - reads without indels (insertions/deletions)
- bowtie2
 - The second version of bowtie2 [Langmead & Salzberg 2012] is optimal for:
 - long reads (more than 50 bp)
 - reads with indels
 - various alignment options
- Each version has its own index file format (bowtie-build / bowtie2-build tools).
- A popular RNA-seq analysis toolset (tophat, cufflinks) is based on bowtie / bowtie2

<http://bowtie-bio.sourceforge.net>

Genome mapping and tools – overview III

- `bwa`
 - `bwa backtrack` [Li, Durbin 2009]:
 - for short reads (< 100bp)
 - `bwa bwsw` [Li, Durbin 2010]:
 - for long reads (70bp - 1Mbp)
 - short indels
 - `bwa mem` [Li 2013]:
 - for long reads (70bp - 1Mbp)
 - faster and more efficient

Genome mapping and tools – overview IV

- samtools package - A set of utilities for processing SAM/BAM files
- samtools view
 - convert a bam file into a sam file - `samtools view sample.bam > sample.sam`
 - Convert a sam file into a bam file - `samtools view -bS sample.sam > sample.bam`
 - Extract all the reads aligned to the range specified. An index of the input file is required
`samtools view -h -b sample_sorted.bam "chr1:10-13" > tiny_sorted.bam`
- samtools sort
`samtools sort unsorted_in.bam sorted_out`
- samtools index
`samtools index sorted.bam (creates an index file, sorted.bam.bai)`
<http://samtools.sourceforge.net>

Genome mapping and tools – overview IV

- samtools package - A set of utilities for processing SAM/BAM files

- samtools view

- convert a bam file into a sam file - `samtools view sample.bam > sample.sam`

- Convert a sam file into a bam file - `samtools view -bS sample.sam > sample.bam`

- Extract all the reads aligned to the range specified. An index of the input file is required

- `samtools view -h -b sample_sorted.bam "chr1:10-13" > tiny_sorted.bam`

- samtools sort

- `samtools sort unsorted_in.bam sorted_out`

- samtools index

- `samtools index sorted.bam (creates an index file, sorted.bam.bai)`

<http://samtools.sourceforge.net>

-b: output BAM
-S: read SAM

add a proper header

Genome mapping and tools – overview IV

- **samtools flagstat** – report basic statistics

```
samtools flagstat sample.bam
```

An example of output:

```
4198456 + 0 in total (QC-passed reads + QC-failed reads)
0 + 0 duplicates
4022089 + 0 mapped (95.80%:-nan%)
4198456 + 0 paired in sequencing
2099228 + 0 read1
2099228 + 0 read2
3796446 + 0 properly paired (90.42%:-nan%)
4013692 + 0 with itself and mate mapped
8397 + 0 singletons (0.20%:-nan%)
167574 + 0 with mate mapped to a different chr
72008 + 0 with mate mapped to a different chr (mapQ>=5)
```

- **samtools faidx** – index a FASTA file

```
samtools faidx ref.fasta (creates an index file ref.fasta.fai)
```
- **samtools merge** – merge several BAM files into one

```
samtools merge out.bam in1.bam in2.bam
```

Genome mapping and tools – overview V

- **BedTools package**
 - `bamtobed` - Convert BAM alignments to BED (& other) formats
 - `bamtofastq` - Convert BAM records to FASTQ records
 - `bedtobam` - Convert intervals to BAM records
 - `closest` - Find the closest, potentially non-overlapping interval
 - `complement` - Extract intervals `_not_` represented by an interval file
 - `coverage` - Compute the coverage over defined intervals
 - `genomecov` - Compute the coverage over an entire genome
 - `getfasta` - Use intervals to extract sequences from a FASTA file
 - `intersect` - Find overlapping intervals in various ways
 - `shuffle` - Randomly redistribute intervals in a genome
 - `sort` - Order the intervals in a file

Genome mapping and tools – overview V

`bedtools intersect`

- Report the intervals that represent overlaps between your two files:

```
bedtools intersect -a cpg.bed -b exons.bed
```

- Report the original feature in each file:

```
bedtools intersect -a cpg.bed -b exons.bed -wa -wb
```

- How many base pairs of overlap were there?

```
bedtools intersect -a cpg.bed -b exons.bed -wo
```

- Counting the number of overlapping features:

```
bedtools intersect -a cpg.bed -b exons.bed -c
```

- Find features that DO NOT overlap:

```
bedtools intersect -a cpg.bed -b exons.bed -v
```



Exercise

- There are paired reads of some DNA sequencing experiment of the human sample:
`bio-bootcamp/R1.fastq.gz`
`bio-bootcamp/R2.fastq.gz`
- You will study some particular region of the human genome
- Map reads to the human reference genome (version hg19 – find path on our Bio Resources)
- Extract reads that map to your region only
- Upload the reads to UCSC genome browser as a custom track
- count the number of insertions and deletions in SAM file

How To: Mapping

- load bowtie2 program:

```
module purge
module load gcc/9.2.0
module load bowtie2
```

- Copy sequence of a chromosome your region is located at as a FASTA file
 - find the path on our website in Homo sapiens > UCSC > hg19 > Chromosomes (7th column)
 - <https://carc.usc.edu/user-information/bio-resources/reference-genomes>
 - Add chr*.fa at the end of the path
 - `cp path_above /home1/ttrojan/bio-bootcamp/results/read-mapping`
- Map reads to this chromosome using `bowtie2` with the standard parameters

Don't forget to make an index (`bowtie-build2`) of the chromosome before mapping!

- You will get a SAM file as an output, convert to BAM (`samtools view`)
- Count the number of insertions and deletions in SAM file (use `cut` for field 6, and `grep`)

How To: Extracting reads

- load BedTools:

```
module purge
```

```
module load gcc/9.2.0
```

```
module load bedtools2
```

- Create a tab-delimited BED file with the coordinates of your region:

```
chr21 10000000 20000000
```

- Convert SAM file with mapped reads to BAM file using `samtools view`
- Use `bedtools intersect` to extract the reads from the BAM file

You'll need a BED file to upload the result to UCSC genome browser, so figure out how to make `bedtools intersect` to produce an output in BED format.

How To: UCSC custom track

- Upload the BED file to UCSC genome browser
'Add custom track' button → Choose file → Submit

Thanks to the whole CARC team:

BD Kim
Bill Jendrzek
Jimi Chu
James K Hong
Chris Taylor
Derek Strong
Cesar Sul
Marco Olguin
Ryan Sim
Carolina Jin

Reference material:

- HPCBio (Holmes J., Clark L., Drnevicch J., Valizadegan N.)
- CNRG (Davidson D., Leigh J.)
- Skoltech (Khrameeva E.)

Mapping exercise: Answers

```
#!/bin/bash
#SBATCH --partition main
#SBATCH --nodes 1
#SBATCH --ntasks 20
#SBATCH --time 01:00:00
#SBATCH --mem 4g
#SBATCH --account=ttrojan_001
#SBATCH --chdir /home1/ttrojan/workshop-bioresources
module purge
module load gcc/9.2.0
module load bowtie2
module load samtools
module load bedtools2
mkdir results/read-mapping
cp data/R*.gz results/read-mapping
gunzip results/read-mapping/R1.fastq.gz
gunzip results/read-mapping/R2.fastq.gz
cp -v /project/biodb/genomes/Homo_sapiens/UCSC/hg19/Sequence/Chromosomes/chr21.fa results/read-mapping/
bowtie2-build --threads $SLURM_NTASKS results/read-mapping/chr21.fa results/read-mapping/chr21index
bowtie2 --threads $SLURM_NTASKS -x results/read-mapping/chr21index -q results/read-mapping/R1.fastq > results/read-mapping/R1.sam
bowtie2 --threads $SLURM_NTASKS -x results/read-mapping/chr21index -q results/read-mapping/R2.fastq > results/read-mapping/R2.sam
samtools view results/read-mapping/R1.bam | cut -f 6 | grep -c 'D' > results/read-mapping/R1.no_of_deletions.txt
samtools view results/read-mapping/R1.bam | cut -f 6 | grep -c 'I' > results/read-mapping/R1.no_of_insertions.txt
samtools view results/read-mapping/R2.bam | cut -f 6 | grep -c 'D' > results/read-mapping/R2.no_of_deletions.txt
samtools view results/read-mapping/R2.bam | cut -f 6 | grep -c 'I' > results/read-mapping/R2.no_of_insertions.txt
samtools view -bS results/read-mapping/R1.sam > results/read-mapping/R1.bam
samtools view -bS results/read-mapping/R2.sam > results/read-mapping/R2.bam
samtools sort results/read-mapping/R1.bam > results/read-mapping/R1_sorted.bam
samtools sort results/read-mapping/R2.bam > results/read-mapping/R2_sorted.bam
samtools view -h -b results/read-mapping/R1_sorted.bam "chr21:100000000-200000000" > results/read-mapping/R1_sorted_region.bam
samtools view -h -b results/read-mapping/R2_sorted.bam "chr21:100000000-200000000" > results/read-mapping/R2_sorted_region.bam
bamToBed -i results/read-mapping/R1_sorted_region.bam > results/read-mapping/R1_sorted_region.bed
bamToBed -i results/read-mapping/R2_sorted_region.bam > results/read-mapping/R2_sorted_region.bed
bedtools intersect -a results/read-mapping/R1_sorted_region.bed -b results/read-mapping/R2_sorted_region.bed > results/read-mapping/reads.bed
```